

# Why Do we need Container Orchestration?

# Containers are Good...


- Both Linux Containers and Docker Containers
  - Isolate the application from the host

FASTER, RELIABLE, EFFICIENT, LIGHTWEIGHT, AND SCALABLE



# Containers Problems!

- Both Linux Containers and Docker Containers
  - Isolate the application from the host




FASTER, RELIABLE, EFFICIENT, LIGHTWEIGHT, AND SCALABLE



# Containers Problems!

- Both Linux Containers and Docker Containers
  - Isolate the application from the host

FASTER, RELIABLE, EFFICIENT, LIGHTWEIGHT, AND SCALABLE



# Problems with Scaling up Containers

It was not Scalable



1

Containers could communicate with each other

2

Containers had to be deployed appropriately

3

Containers had to be managed carefully

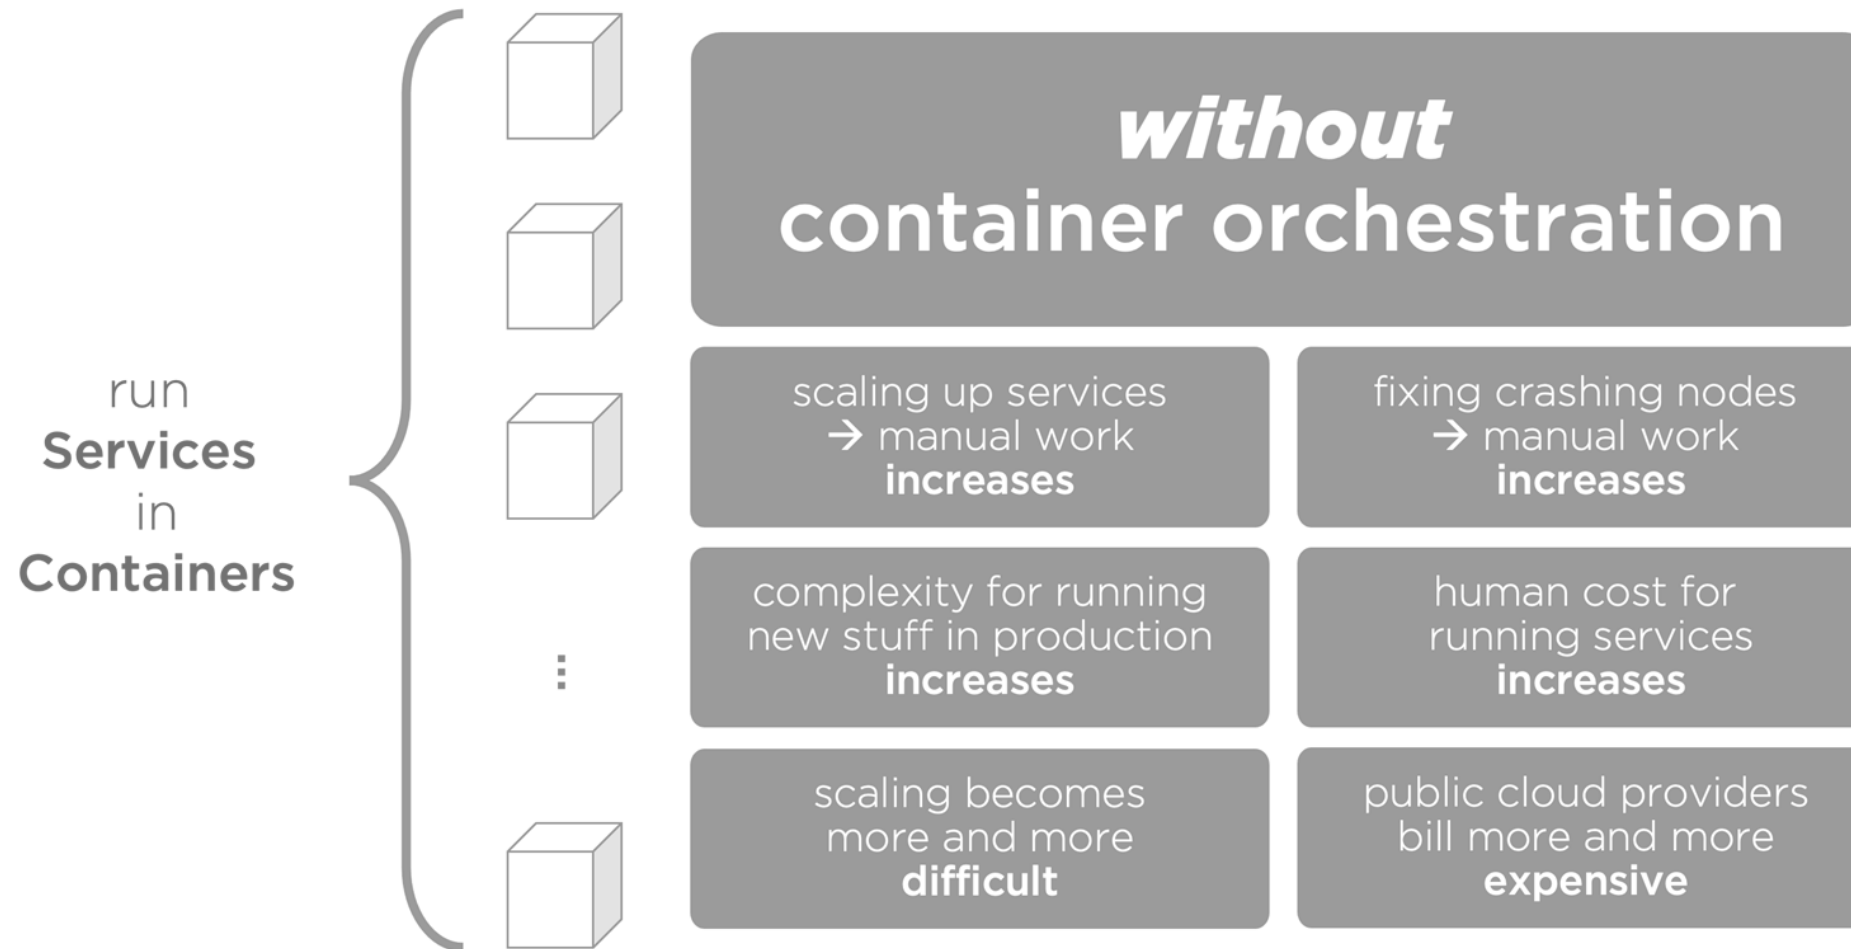
4

Auto scaling was not possible

5

Distributing traffic was still challenging

# Containers Without Orchestration



# Container Orchestration


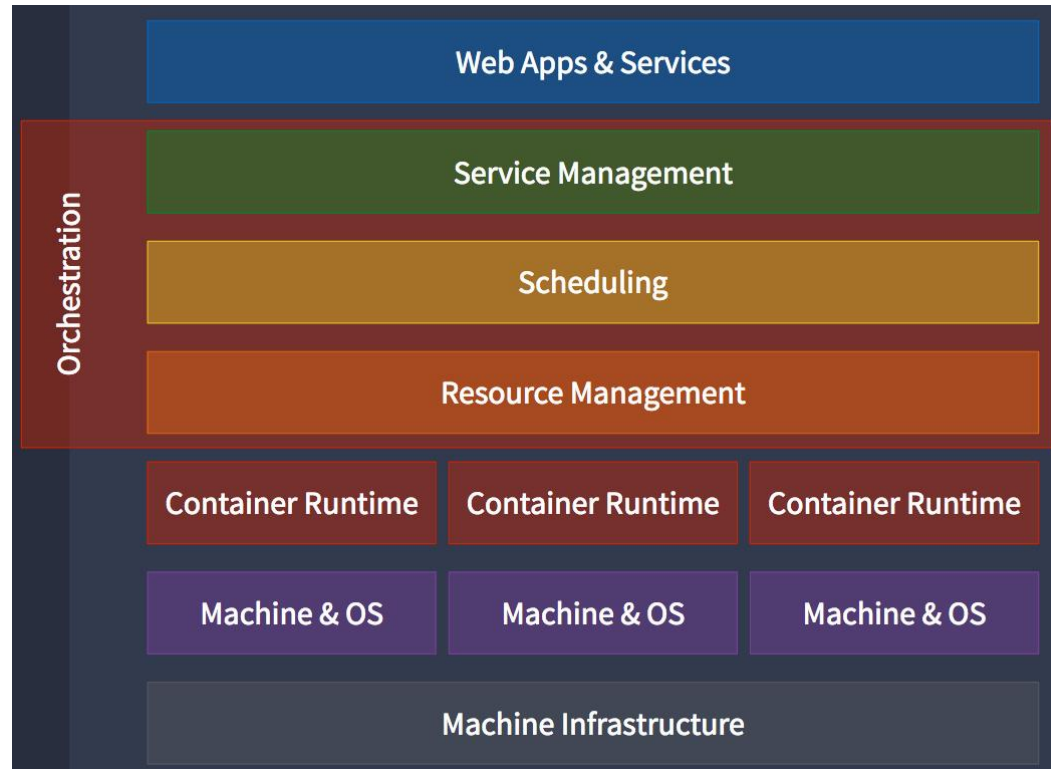
# What is container orchestration?

- Container orchestration is the automation of much of the operational effort required to run containerized workloads and services.
- This includes a wide range of things software teams need to manage a container's lifecycle, including:
  - provisioning,
  - deployment,
  - scaling (up and down),
  - networking,
  - load balancing and more.



# Container Orchestration

- Container orchestration automates and simplifies provisioning, and deployment and management of containerized applications.



# Container Orchestration


- Container orchestration is the automatic process of managing or scheduling the work of individual containers for applications based on microservices within multiple clusters.
- The widely deployed container orchestration platforms are based on open-source versions like Kubernetes, Docker Swarm or the commercial version from Red Hat OpenShift.

Container Orchestration Software  
(Docker, Openshift & Kubernetes)



**Automate:**

- Configuration
- Provisioning
- Availability
- Scaling
- Security
- Resource allocation
- Load balancing
- Health monitoring



# Container Orchestration


- Fault-tolerance
- On-demand scalability
- Optimal resource usage
- Auto-discovery to automatically discover and communicate with each other
- Accessibility from the outside world
- Seamless updates/rollbacks without any downtime.

# Why Do We Need Container Orchestration?

- Container orchestration is used to automate the following tasks at scale:
  - Configuring and scheduling of containers
  - Provisioning and deployments of containers
  - Availability of containers
  - The configuration of applications in terms of the containers that they run in
  - Scaling of containers to equally balance application workloads across infrastructure
  - Allocation of resources between containers
  - Load balancing, traffic routing and service discovery of containers
  - Health monitoring of containers
  - Securing the interactions between containers.

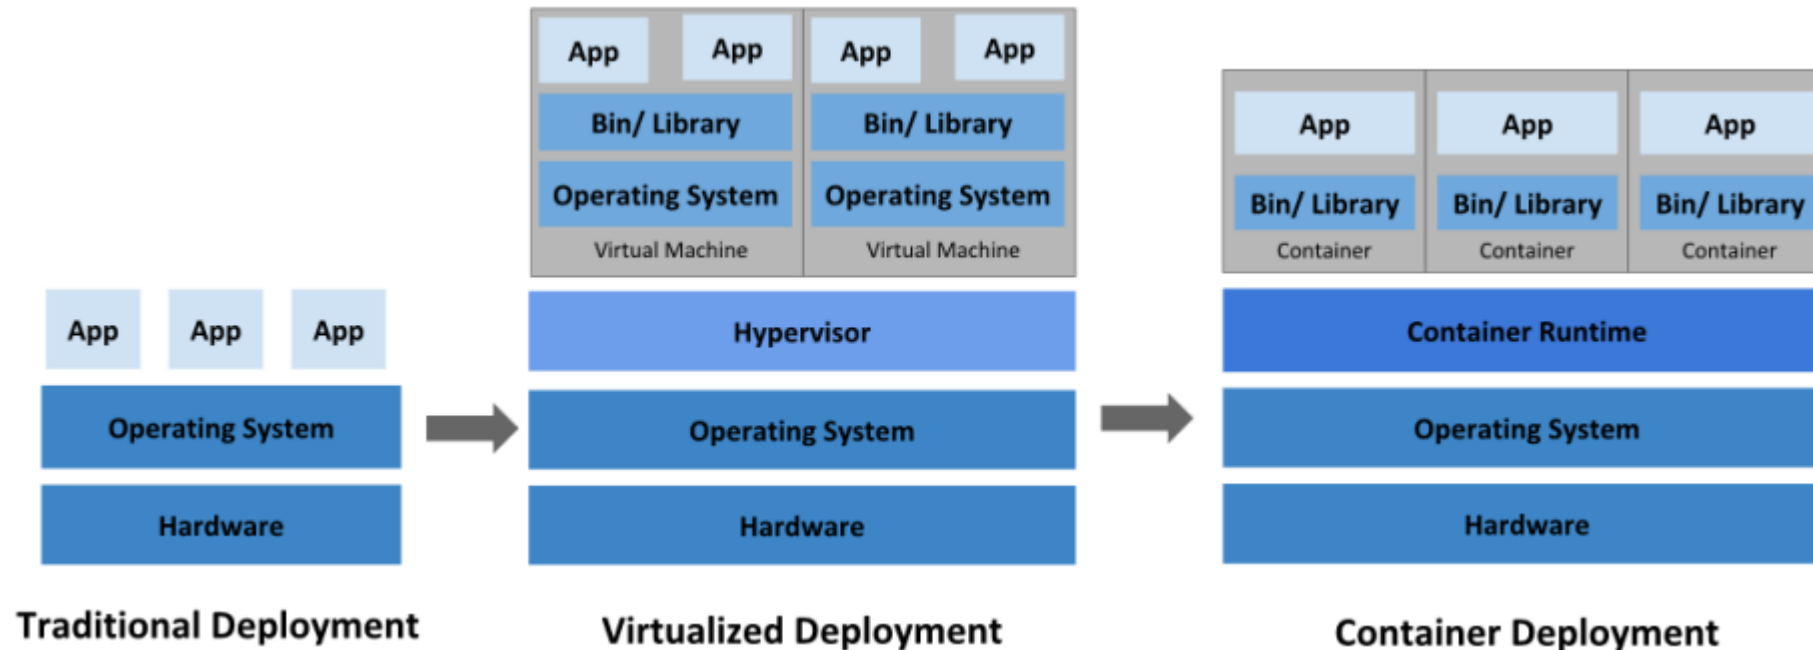
# Cloud Orchestration

- Infrastructure-as-a-Service (IaaS)
  - Provisioning virtual machines from a cloud service provider (CSP)



# Container Orchestration

- Application containers
  - Lightweight OS-virtualization
  - Application packaging for portable, reusable software



# Container Orchestration

Container Orchestration Software  
(Docker, Openshift & Kubernetes)



OPENSIFT




**Automate:**

- Configuration
- Provisioning
- Availability
- Scaling
- Security
- Resource allocation
- Load balancing
- Health monitoring




Application Environment  
w/ Multiple Containers



# Container Orchestration Tools



# Container Orchestration Tools



# What is Kubernetes and How it works

# Kubernetes




Kubernetes is an open-source container management tool which automates container deployment, container (de)scaling and container load balancing

- Benefits (Works with all cloud vendors (Public, Private (on-premises), and Hybrid))

## More about Kubernetes

- Developed by Google and written in Golang with a huge community
- Can group 'n' containers into one logical unit for managing and deploying them




# Kubernetes

- Kubernetes is an open source container orchestration engine for automating deployment, scaling, and management of containerized application.
- Originally an open source project launched by Google and now part of the Cloud Native Computing Foundation (CNCF).
- Kubernetes is highly **extensible** and **portable**
- Kubernetes is considered **highly declarative**
- Kubernetes initial release (**7 June 2014**)
- Releases every 3 months

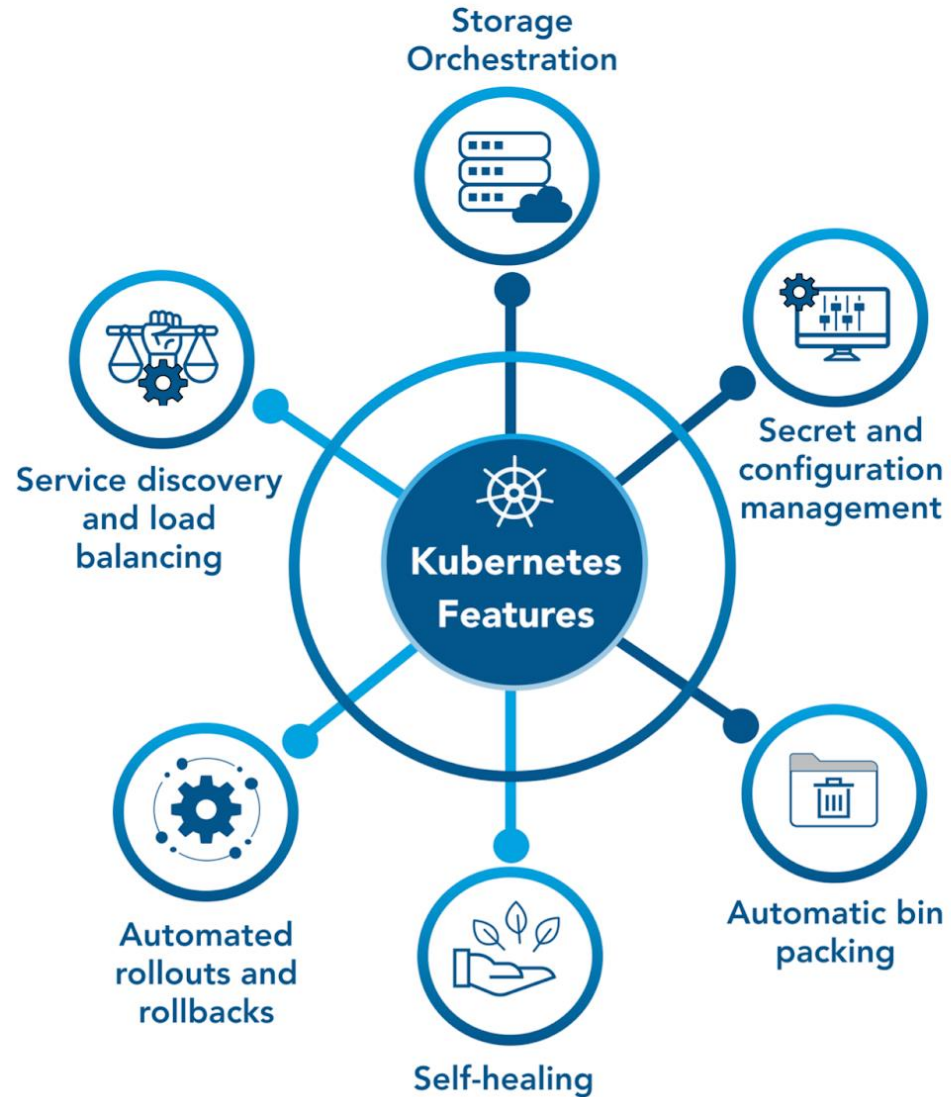


**kubernetes**

# Kubernetes




# Kubernetes Features




# Kubernetes Myths

- Kubernetes is not:
  - To be compared with Docker
  - For containerizing apps
  - For apps with simple architecture
- Kubernetes is actually:
  - Robust and reliable
  - A container orchestration platform
  - A solution for scaling up Containers
  - Backed by huge community



# Kubernetes vs Docker Swarm



vs



kubernetes

## Docker Swarm

1

No Auto Scaling

2

Good community

3

Easy to start a cluster

4

Limited to the Docker API's capabilities

5

Does not have as much experience with production deployments at scale

## Kubernetes

1

Auto Scaling

2

Great active community

3

Difficult to start a cluster

4

Can overcome constraints of Docker and Docker API

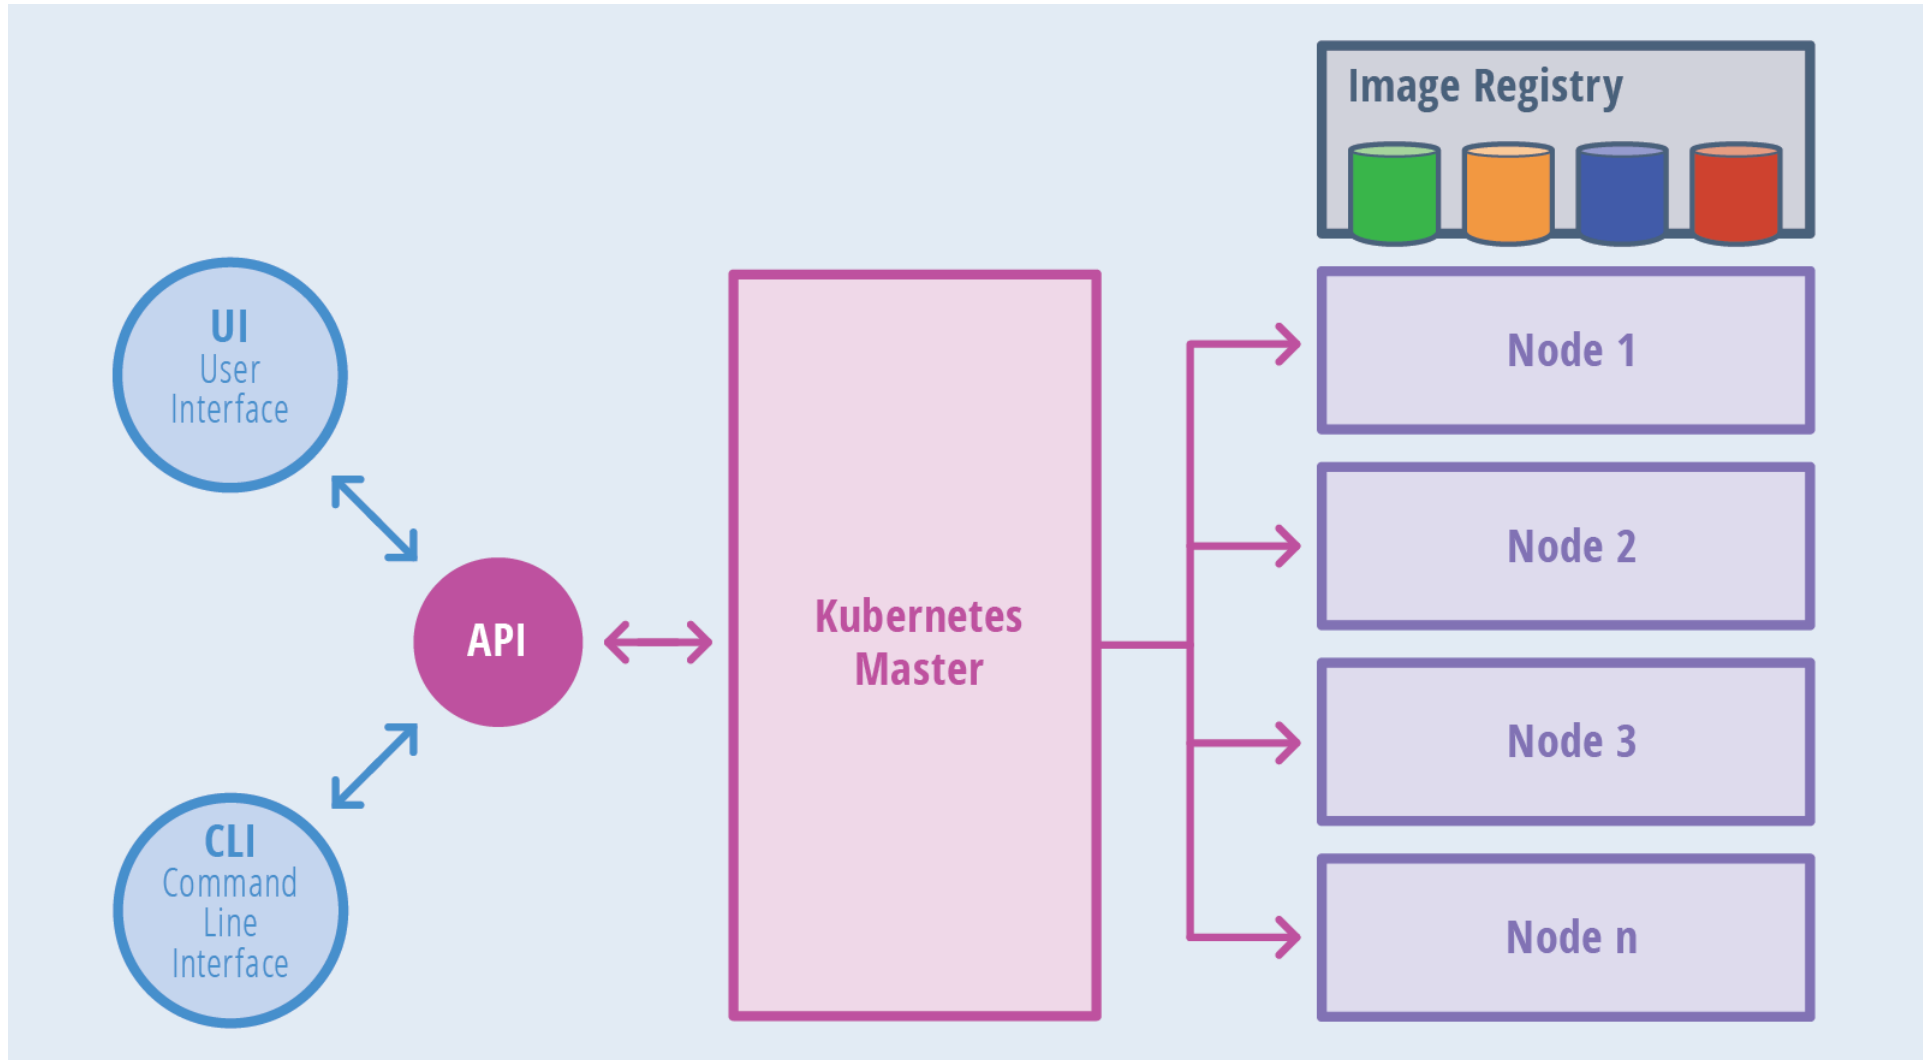
5

Deployed at scale more often among organizations



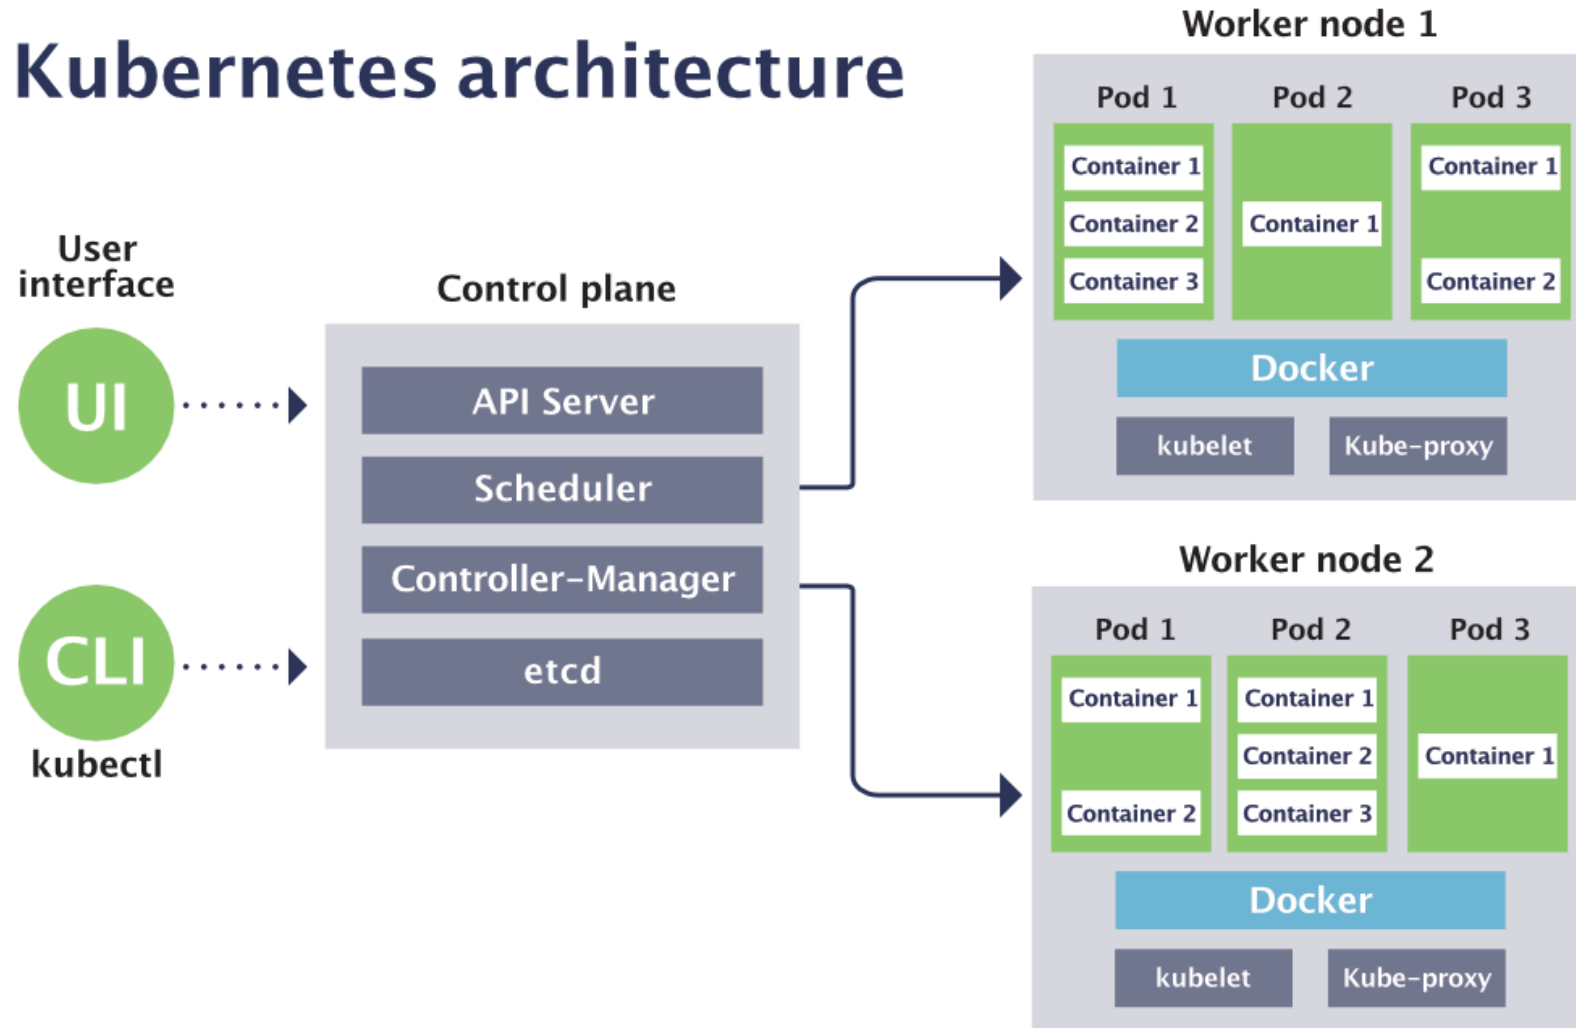
# Kubernetes Architecture

# Kubernetes Architecture

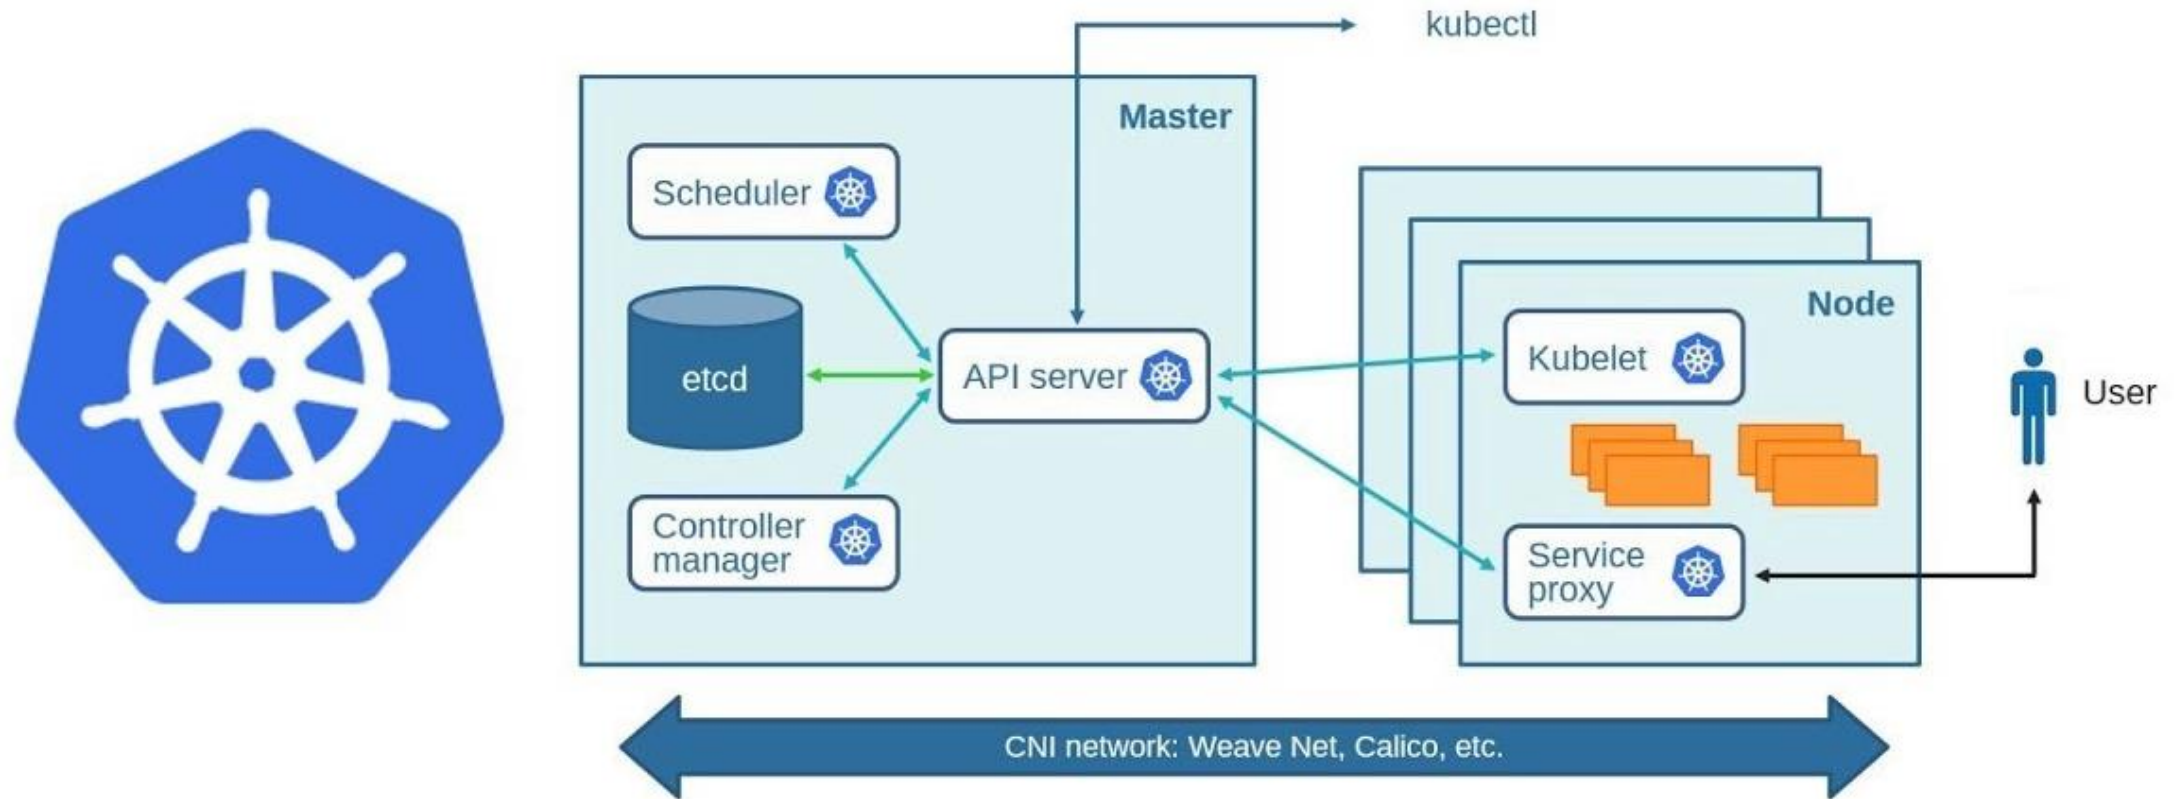


# Kubernetes Architecture

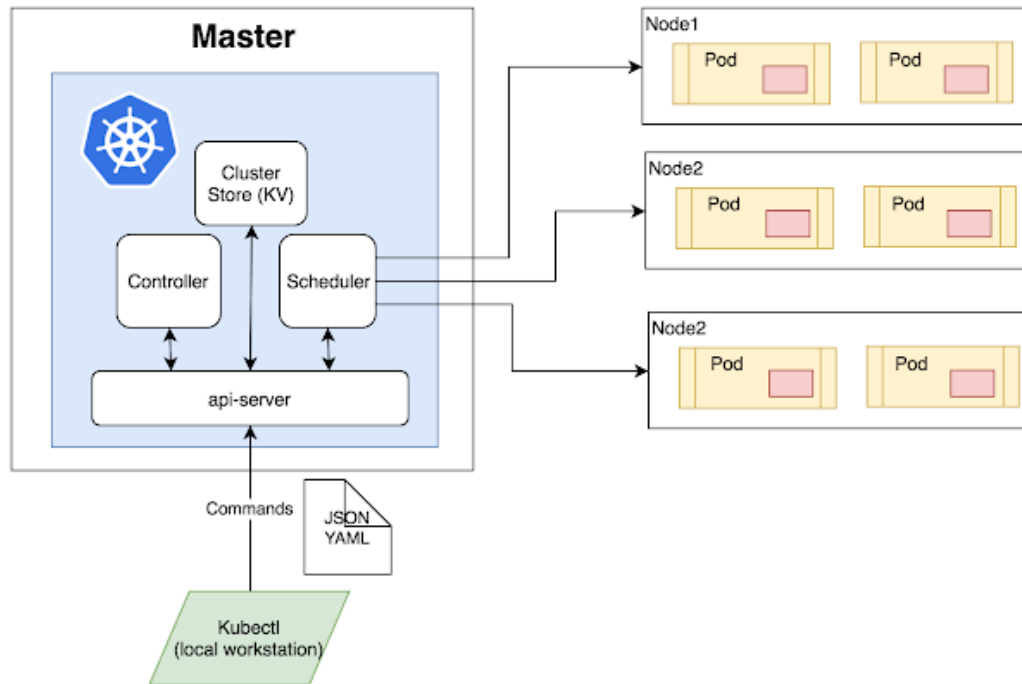
## Kubernetes architecture



# Kubernetes Architecture



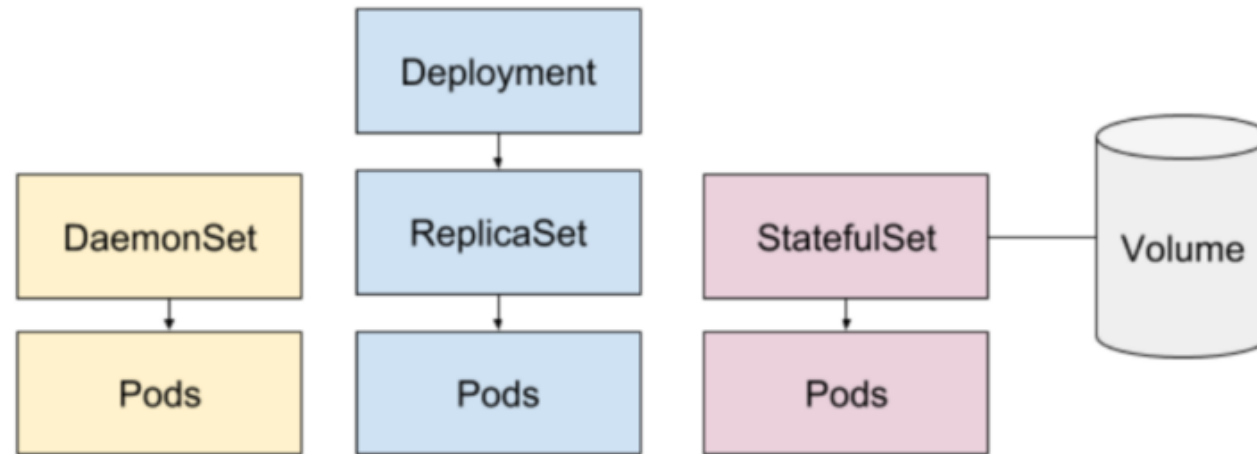
# Kubernetes Components



- **Cluster:**
  - A collection of hosts(servers) that aggregates their available resources.
- **Master:**
  - A collection of components which make up the control panel of Kubernetes.
- **Node:**
  - A single host which is capable of running on a physical or virtual machine.
- **Namespace:**
  - A logical cluster or environment.


# Basic Kubernetes Objects

- Pod
- Deployment, ReplicaSet
- DaemonSet
- StatefulSet
- Service
- Secret



# Pods

- Pods are the smallest deployable units of computing that you can create and manage in Kubernetes.
- A Pod is a group of one or more containers, with shared storage/network resources, and a specification for how to run the containers.
- A Pod's contents are always co-located and co-scheduled.
- To create and manage multiple Pods, Kubernetes defines multiple resource types:
  - Deployment
  - StatefulSet
  - DaemonSet



# Deployment

- Deployment: represents a set of multiple, identical Pods with no unique identities.
  - It runs multiple replicas Pods and automatically replaces any instances that fail or become unresponsive.
  - Deployments ensure that one or more instances of Pods are available to serve user requests.
- PodTemplates are specifications for creating Pods, and are included in resource objects such as Deployment object.
- A ReplicaSet is the next-generation of ReplicationControllers
  - It ensures that a specified number of pod replicas are running at any given time.



# Pod Management

- StatefulSet: manages deployment and scaling of a set of Pods, with durable storage and persistent identifiers for each pod.
  - Unlike a Deployment, a StatefulSet maintains a sticky identity for each of its Pods.
- DaemonSet: ensures that all (or some) nodes run a copy of a Pod.
  - As nodes are added to the cluster, Pods are added to them.
  - As nodes are removed from the cluster, those Pods are garbage collected.
  - Deleting a DaemonSet will clean up the Pods it created.

# Controller

- Controllers are control loops that watch the state of the cluster, then make or request changes where needed.
- A controller tracks at least one Kubernetes resource object.
- The controller(s) for that resource are responsible for making the current state come closer to that desired state (specified in the spec field).
- Kubernetes comes with a set of controllers that run inside the kubecontroller-manager
- The Deployment controller is an example of controller that come as part of Kubernetes itself ("built-in" controllers).

# Service

- The services model relies upon the most basic, though most important, aspect of services: **discovery**.
- a Service is an abstraction which defines a logical set of Pods and a policy by which to access them.
- The set of Pods targeted by a Service is usually determined by a **selector**.
- Services ensure that traffic is always routed to the appropriate Pod within the cluster, regardless of the node on which it is scheduled.
- Each service exposes an **IP address**, and may also expose a **DNS endpoint**, both of which will never change.
  - Internal or external consumers that need to communicate with a set of pods will use the service's IP address, or its more generally known DNS endpoint.

# Secret

- Secrets are secure objects which store sensitive data, such as passwords, OAuth tokens, and SSH keys, in your clusters.
- Storing sensitive data in Secrets is more secure than plaintext in Pod specifications.
- Using Secrets gives you control over how sensitive data is used, and reduces the risk of exposing the data to unauthorized users.

# Kubernetes Cluster

